



Universität  
Rostock



Traditio et Innovatio

# Distributed Algorithms

## Clocks

Univ.-Prof. Dr.-Ing. habil. Gero Mühl

Architecture of Application Systems

Faculty for Computer Science and Electrical Engineering

University of Rostock



# Overview

- > Time and Clocks
- > Synchronization of physical clocks
- > Order of events
- > Logical clocks
  - > Lamport clocks
  - > Vector clocks
  - > Application of vector clocks (causal multicast)

# Time in Distributed Systems

# The Importance of Time

- > Determination of time and the measurement of time durations is indispensable for the coordination of human activities
- > We have internalized the existence of a global time
- > Therefore, our clocks have to be synchronized
  - > Synchronization of clocks by means of church clock, telegraphy, radio, GPS
  - > Clock synchronization made longitude determination in seafaring feasible



© NMM London, MoD Art Collection

# Clocks

- > Model: a **clock** maps the real time  $t$  to a time stamp  $C(t)$
- > **Resolution**
  - > Smallest period of time by that two values of the clock can differ (e.g., 10 ms) → **Tick duration**
- > **Drift**
  - > Deviation of the speed of the clock from real time (ca.  $\pm 10^{-6}$  with quartz clocks, ca.  $\pm 2$ sec per month)
- > **Offset**
  - > Deviation of the clock from real time at a point in time, i.e.,  $t - C(t)$
- > A **perfect clock** has a drift and an offset of 0:  $C(t) = t$




# Correct Clocks

- > A **correct clock** has a *limited maximal drift*  $\rho$

$$(1 + \rho)^{-1} \leq \frac{dC}{dt} \leq 1 + \rho$$

„Speed“ of clock



too fast  
 $dC / dt > 1 + \rho$

exact  
 $dC / dt = 1$

too slow  
 $dC / dt < (1 + \rho)^{-1}$

# Applications of Time Stamps

- > Evaluate actuality of data
- > Performance measurements
- > Determine validity of access authorization
- > Derive total order of events  
(e.g., for synchronization, debugging and audit)
- > Evaluate sensor data and control actuators  
(real-time systems)
- > ...

# Synchronization of Physical Clocks



# Time in Distributed Systems

- > Each computer has its own inaccurate digital clock
- > The drifts of the clocks are different from each other
- > Without synchronization, the values of the clocks can differ arbitrarily from each other → **clock synchronization**
- > Clock synchronization in distributed systems is only possible through message exchange
- > Hereby, the message delay plays an important role

# Synchronization of Correct Clocks

- > Two correct clocks with drift  $\rho$  should *not* deviate by more than time  $d$
- > How long is the maximum possible interval for synchronization?
- > Assumptions
  - > At  $t = 0$ , the clocks are synchronized:  $C_1(0) = C_2(0)$
  - > Worst case: One clock is as fast as drift allows, the other as slow as drift allows
  - > Without loss of generality:
    - >  $C_1(t) / t = 1 + \rho$
    - >  $C_2(t) / t = 1 / (1 + \rho)$

# Synchronization of Correct Clocks

> Therefore

$$(1 + \rho) t - \frac{t}{1 + \rho} \leq d$$

$$\frac{(1 + \rho)^2 t - t}{1 + \rho} \leq d$$

$$t \frac{2\rho + \rho^2}{1 + \rho} \leq d$$

$$t \leq d \frac{1 + \rho}{2\rho + \rho^2}$$

> Clocks must be synchronized again before  $d \frac{1 + \rho}{2\rho + \rho^2}$

> For very small  $\rho$ : synchronization before  $\rho^{-1} d / 2$

# Synchronization Against a Perfect Clock

- > Assumption: Clocks are synchronized at  $t = 0$
- > Two cases:

Clock as fast as possible

$$(1 + \rho) t - t \leq d$$

$$\rho t \leq d$$

$$t \leq \frac{d}{\rho}$$

Clock as slow as possible

$$t - \frac{t}{1 + \rho} \leq d$$

$$\frac{(1 + \rho)t - t}{1 + \rho} \leq d$$


$$\frac{\rho t}{1 + \rho} \leq d$$

$$t \leq d \frac{1 + \rho}{\rho}$$

- $\frac{d}{\rho} \leq d \frac{1 + \rho}{\rho}$
- > Since  $\frac{d}{\rho} \leq d \frac{1 + \rho}{\rho}$ , the clock must be synchronized again before  $\rho^{-1} d$

# External Clock Synchronization [5]


- > Scenario: Process  $P_2$  wants to adjust its clock on  $P_1$
- > Assumptions
  - >  $P_1$  has a clock with an assumed drift of 0
  - > Message delay  $x$  is known
- >  $P_2$  adjusts its clock by  $t_2 + x - t_3$
- > Error of adjustment  $e_{max} = 0$
- > If the maximal deviation shall stay smaller than  $d$ ,  
a new synchronization is at the latest necessary after  $\rho^{-1} d$



→  $P_2$  sets its clock  
backwards by 1


# Unforeseeable Message Delay

- > In reality, message delays are (in nearly all cases) load-dependent and unbounded
  - > delay is higher with a high load than with a low load
  - > delay can be arbitrarily long
- > Then, the preceding procedure leads only to an approximate adjustment




# Synchronization with Unforeseeable Delays

- > Round Trip Time (RTT)  $D := (t_3 - t_1)$
- > Let  $t$  be the local time of  $P_1$ , when  $P_2$  has the local time  $t_3$
- >  $t$  lies in the interval  $[t_2 + \text{min}, t_2 + D - \text{min}]$
- > Without further knowledge, the best prediction possible for  $t$  is the middle of the interval, i.e.,  $t_2 + D / 2$
- > Thus,  $P_2$  corrects its clock by  $t_2 + D / 2 - t_3$
- > The maximal adjustment error is  $e_{\max} = D / 2 - \text{min}$



# Probabilistic Limitation of the max. Error

- > Idea: 1. Only accept the value of the reference clock if  $D \leq U$  for a **given bound**  $U > 2 \cdot \min$   
2. If that fails, repeat the attempt at most  $k$ -times, always after a **waiting time**  $W$
- > Let  $p$  be the probability that  $D \geq U$  for one attempt
- > Probability for  $k$  failures after another is then  $q = p^k$
- > Expected number of attempts until success is  $E = 1 / (1 - p)$






# Determination of the Synchronization Interval

- > The synchronization interval depends on the error of the last *successful* attempt
- > To keep the deviation below  $d$ , the next success after a successful synchronization with error  $e$  must be after at most  $\Delta = \rho^{-1} (d - e)$  time
  - > Minimal synchronization interval  $\Delta_{min} = \rho^{-1} (d - e_{max})$
  - > Maximal synchronization interval  $\Delta_{max} = \rho^{-1} d$
- > First, the local clock is adjusted; then, the next attempt is started after  $\Delta - k W$  at the latest
- > Minimal maximal deviation with immediate start ( $\Delta = k W$ ) is
$$d_{min} = k W \rho + e_{max}$$

# Adjustment of the Local Clock Time


- > Requirements
  - > Great leaps of clock time must be avoided
  - > Clock time must not decrease
- > Solution
  - > Local clock is run slower or faster, until the offset is fully compensated



# Order of Events

# Order of Events in Distributed Systems

- > Send and receive of a message are events  
→ send event, receive event
- > Additionally, there are local events



- > In distributed systems, the absolute time of events are often not important; it often suffices to order the events

# Order Relations

- > An order relation  $<$  is a

- > *irreflexive* for no event  $a$  applies  $a < a$

- > *asymmetrical*  $a < b \Rightarrow \neg(b < a)$

- > *transitive*  $a < b \wedge b < c \Rightarrow a < c$

binary relation on the set of all events  $E$

- > **Partial order:**

The order relation *is not* defined for *all* pairs of events

- > **Total order:**

The order relation is defined for *all* pairs of events, i.e.,

$$e_1 \neq e_2 \Rightarrow e_1 < e_2 \vee e_2 < e_1$$

# Possible Order Requirements

- > FIFO (first in first out) order
- > Causal order
- > Total delivery order
- > FIFO-total order
- > Causal-total order

# FIFO Order

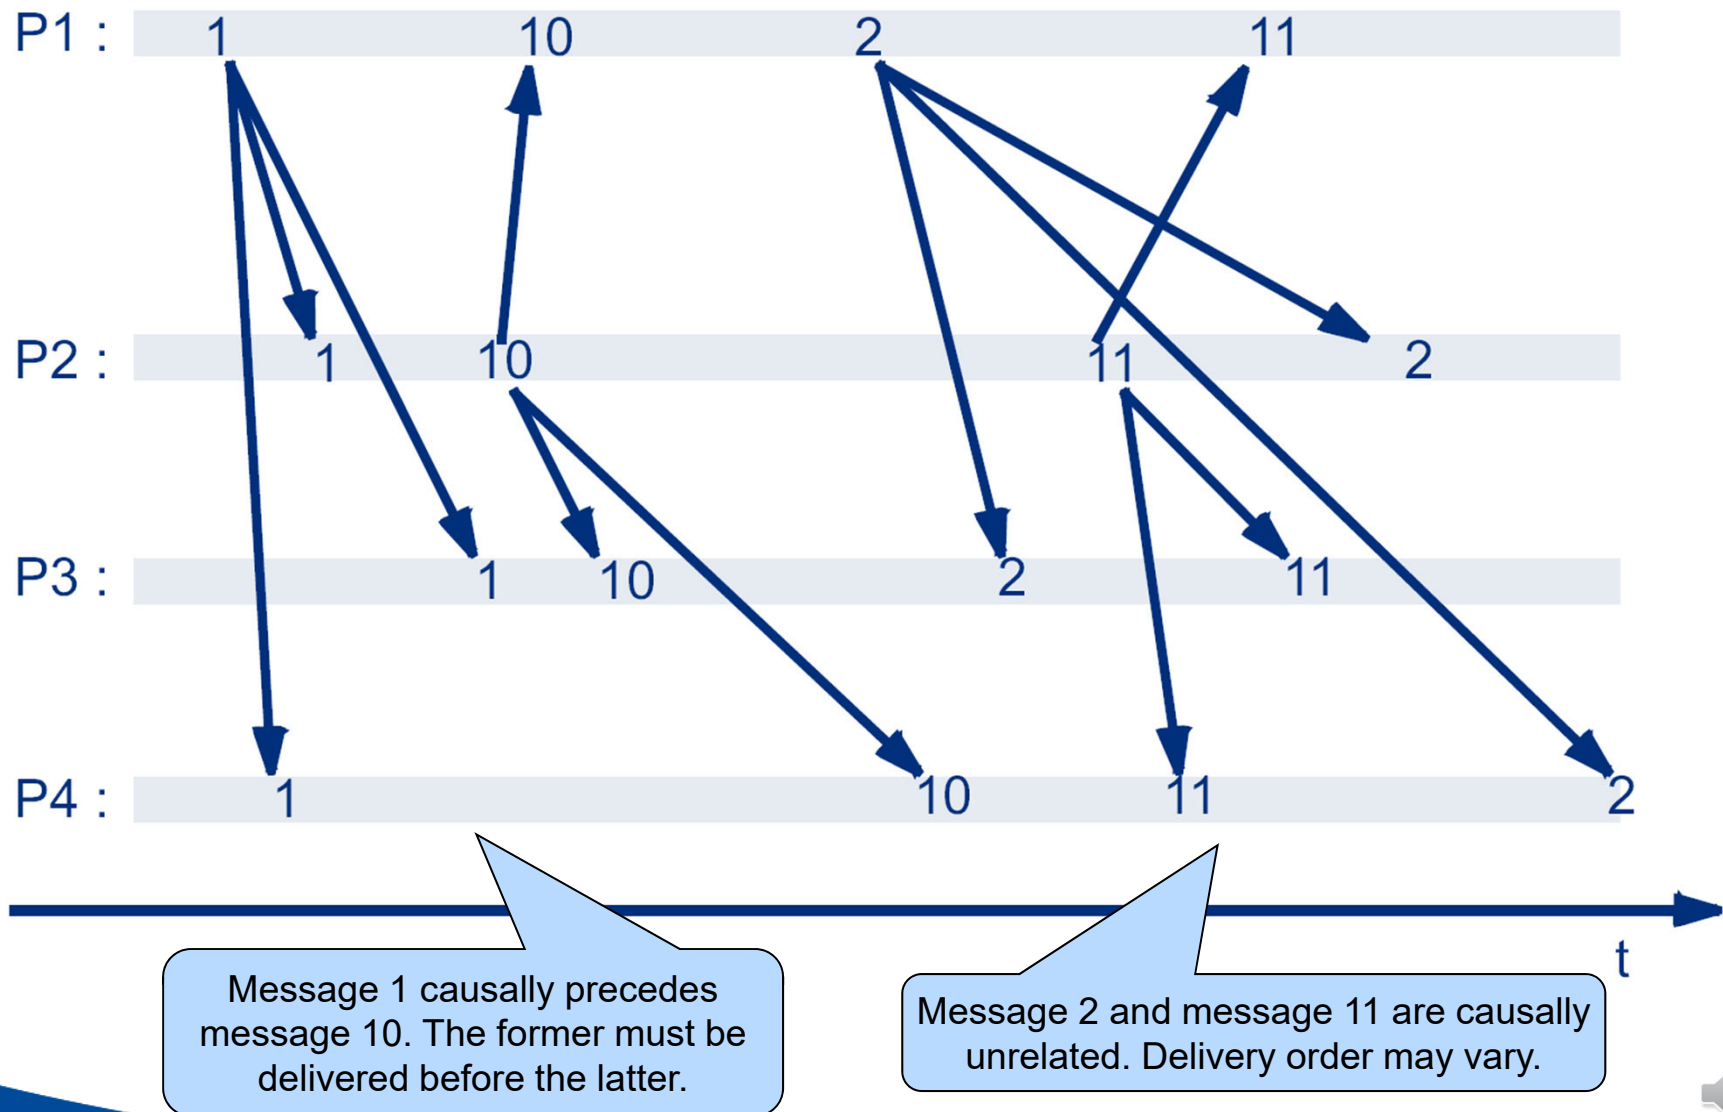
- > If a process sends message  $m_1$  before message  $m_2$ ,  $m_1$  is delivered before  $m_2$  to the receiver(s)
- > Since messages can overtake each other, we distinguish between “receive” and “deliver”
- > Messages are, then, delayed if necessary to enforce an order

# Causal Order

- > If the sending of message  $m_2$  causally depends on the sending of the message  $m_1$ , then  $m_1$  is delivered before  $m_2$  to any receiver (getting both messages)
- > Causal order is *stronger* than FIFO
  - > Every causal order is also FIFO ordered
  - > The inverse does not apply
- > Causality is *hypothetical*
  - > Every possible causality is preserved
  - > It is, thus, possible that causally ordered events do not depend on each other in reality




# Example: Causally Ordered Multicast



# Total Delivery Order

- > Only useful, if a message is delivered to more than one process → **multicast communication**
- > If two processes  $P$  and  $Q$  both deliver the messages  $m_1$  and  $m_2$ , then  $P$  delivers  $m_1$  before  $m_2$  if and only if  $Q$  also does.
- > This means that all processes receiving  $m_1$  and  $m_2$ , deliver the messages in the same order
- > Caution: The order itself is not specified, especially, causality can be violated
- > A multicast with total delivery order is also called **atomic multicast**
- > Caution: A total delivery order is not always total in the sense of a total order relation (slide 29)

# Example for Total Delivery Order



# FIFO-Total and Causal-Total Order

- > FIFO-Total order


- > An order that is both FIFO and total with respect to delivery

- > Causal-Total order

- > An order that is both causal and total with respect to delivery

- > As causal order implies FIFO order, a causal-total order is also a FIFO-total order

# Example for FIFO-Total Order



# Logical Clocks

# A Simple Logical Time

- > A **logical clock** assigns a time stamp  $C(e)$  to each event  $e$
- > Each process  $P_i$  manages a counter  $C_i$  that is increased by 1 when an event  $e$  occurs
- > The event  $e$  gets the new value of the counter as **logical time stamp**
- > The logical time stamps  $C(e)$  define a partial order on the set of events  $e_1 < e_2 \Leftrightarrow C(e_1) < C(e_2)$
- > It is partial as some events might get the same timestamp

# A Simple Logical Time

- > It is possible to supplement the simple logical time to a total order through the usage of unique process IDs as tiebreaker
- > The time stamp  $C'(e_i)$  of an event  $e_i$  is a pair  $(C_i, P_i)$ 
$$e_1 < e_2 \Leftrightarrow C'(e_1) < C'(e_2)$$
$$\Leftrightarrow C_1 < C_2 \vee C_1 = C_2 \wedge P_1 < P_2$$
- > Since process identities are unique, for two arbitrary events
$$e_1 \neq e_2 \Rightarrow e_1 < e_2 \vee e_2 < e_1$$



# A Simple Logical Time – Problem

- > Problem: The simple logical time does not take the causal relation between events into account
- > Example: Replicated article data base
  - > Editor adds a new article
  - > Chief editor redacts the article afterwards
  - > If the second action gets a smaller time stamp than the first one, the actions are applied in the wrong order and the data base contains the article which was not redacted

# Happened-Before Relation (Lamport, 1978)

> The relation  $\rightarrow$  („happened before“) on the set of events is the *smallest* order relation that fulfills the following 3 conditions

1. If  $a$  and  $b$  are two events in a process and  $a$  occurs before  $b$ , then  $a \rightarrow b$
2. If  $a$  is the sending of a message in a process and  $b$  the receipt of the same message in another process, then  $a \rightarrow b$
3. If  $a \rightarrow b$  and  $b \rightarrow c$ , then  $a \rightarrow c$

“Smallest” here means that the relation contains exactly those pairs satisfying the conditions.


Condition 3 is implicitly implied by requiring that the relation is an order.

> An event  $b \neq a$  **causally depends** on  $a$ , if  $a \rightarrow b$

> Two events  $a \neq b$  are **causally independent**, written  $a \parallel b$ , if neither  $a \rightarrow b$  nor  $b \rightarrow a$

# Happened-Before Relation – Interpretation

- >  $a \rightarrow b \Rightarrow$  “ $b$  causally depends on  $a$ ”
- >  $a \parallel b \Rightarrow$  „ $a$  and  $b$  have *not* influenced each other causally“
- >  $a \rightarrow b \Leftrightarrow$  „One can get from  $a$  to  $b$  in the space-time diagram by following the process lines and message lines from  $a$  in the direction of *ascending* time“
- > For the example below applies, e.g.,  $e_{1,1} \rightarrow e_{3,4}$  and  $e_{1,6} \parallel e_{2,6}$



# Clock Condition (Lamport, 1978)

- > Requirement for a logical clock (**clock condition**)
  - > For all events  $a, b$ :  $a \rightarrow b \Rightarrow C(a) < C(b)$
  - > Thus, the clock *preserves* the causal order of the events
- > Attention: implication only holds in one direction!
  - > It only applies  $C(a) < C(b) \Rightarrow a \rightarrow b \vee a \parallel b$
- > From the clock condition follows:  $C(a) = C(b) \Rightarrow a \parallel b$
- > How can a logical clock be realized that fulfills the clock condition?


# Lamport's Clocks – Realization

- > Each process  $P_i$  has a local logical clock  $L_i$ , whose value is adapted when an event occurs
- > A local event occurs at process  $P_i$ 
  - > Value of the local clock  $L_i$  is increased by one
  - > Event gets the new value as time stamp
- >  $P_i$  sends a message
  - > Value of local clock  $L_i$  is increased by one
  - > Send event gets new value as time stamp
  - > Message sent carries time stamp of its send event  $t_m$

# Lamport's Clocks – Realization

- >  $P_i$  receives a message  $m$  with time stamp  $t_m$ 
  - > Value of the local clock is adapted to
$$L_i := \max(L_i, t_m) + 1$$
  - > Receive event gets the new value as time stamp
- > Incrementing after maximum calculation ensures that the receive event gets a higher time stamp than both the send event and the preceding local event at the receiver

# Lamport's Clocks – Example



# Lamport's Clocks – Characteristics

- > Lamport's clocks fulfill the clock condition!
- > The logical time stamps  $L(e)$ , thus, define a *partial* order on the set of events maintaining the causal relation
- > Complementing it to a *total* order is possible again
- > Problem
  - > By means of a time stamp one can not tell for sure, whether two events are causally related
  - > For that purpose also the converse of the clock condition would have to apply



# Vector Clocks – Motivation

- > Assume we had a clock that turns the implication of the clock condition into an equivalence
  - >  $a \rightarrow b \Leftrightarrow C(a) < C(b)$
- > With such a clock, we could determine how events are related by means of the time stamps of the events
  - >  $C(a) < C(b) \Rightarrow a \rightarrow b$
  - >  $C(b) < C(a) \Rightarrow b \rightarrow a$
  - >  $\neg (C(a) < C(b) \vee C(b) < C(a)) \Rightarrow a \parallel b$
- > For all three equations also the converse applies
- > For  $\neg (C(a) < C(b) \vee C(b) < C(a))$ , we write  $C(a) \parallel C(b)$
- > How can a clock with such characteristics be realized?

# Vector Clocks – (Mattern, Fidge, 1988)

- > Each process  $P_i$  holds a vector time stamp  $V_i$  consisting of  $n$  counters that are initially all zero
- > Local event
  - > If an event occurs in a process  $P_i$ , it increments the  $i$ -th component of its vector
- > Sending a message
  - > If  $P_i$  sends a message, the new version of  $V_i$  is sent along with the message
- > Receiving a message
  - > If  $P_i$  receives a message with vector time stamp  $T$ , it assigns to  $V_i$  the component-wise maximum of the  $T$  and the new version of  $V_i$
- > The vector time can (also) be extended to a *total* order by using process identities as tiebreakers

# Vector Clocks

- > Component-wise maximum of two vectors

$$\max(V_i, V_j) := (\max(V_i[1], V_j[1]), \dots, \max(V_i[n], V_j[n]))$$

- > Component-by-component comparison of two vectors

$$V_i < V_j \Leftrightarrow V_i \neq V_j \wedge \forall 1 \leq k \leq n : V_i[k] \leq V_j[k]$$

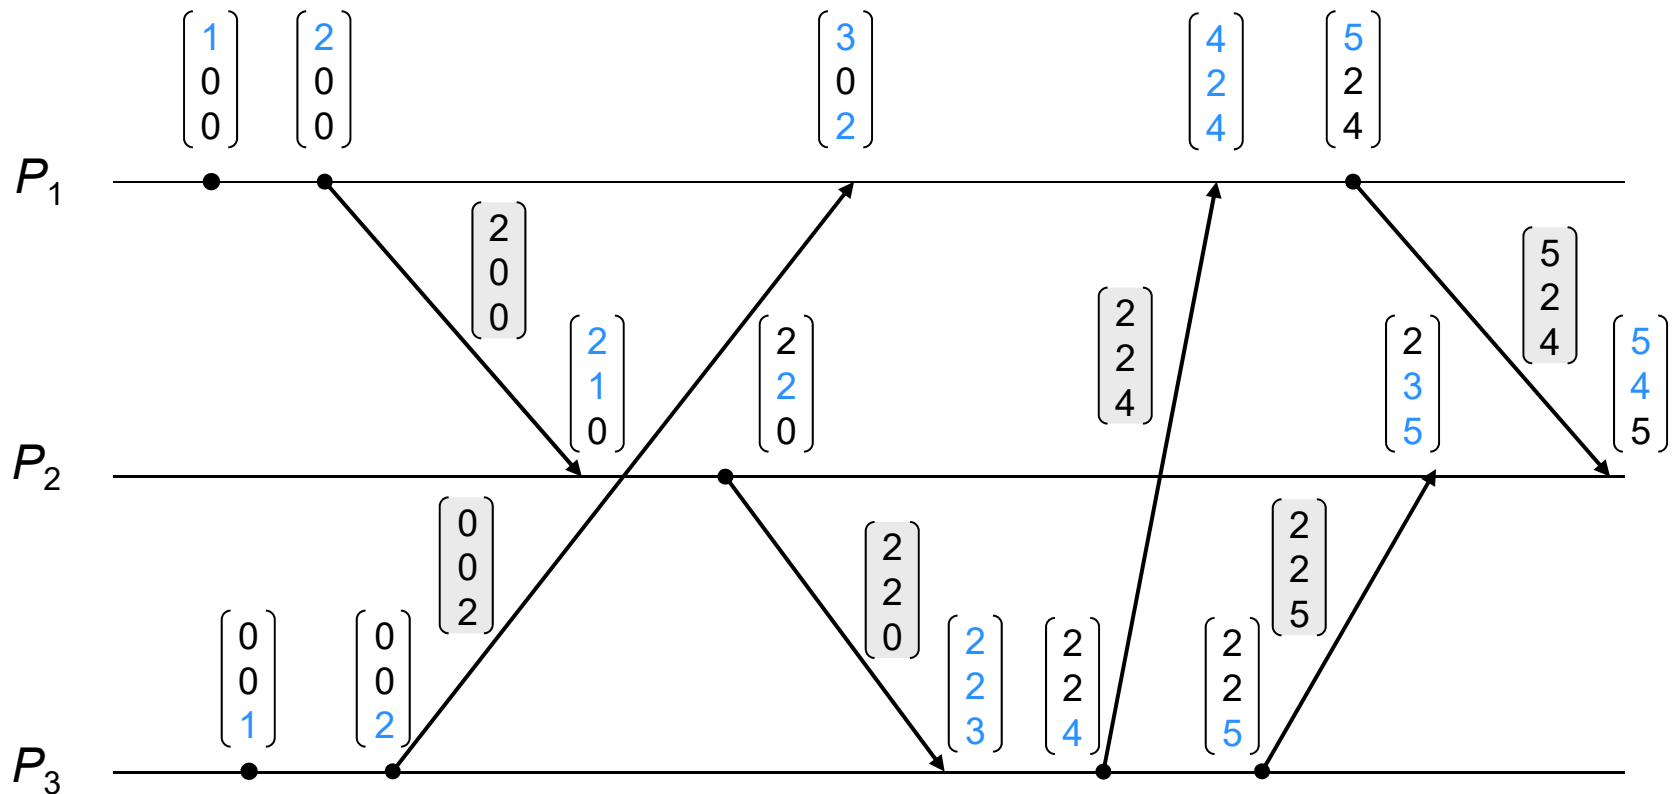
- > Vector clocks only define a *partial* order!

$$\begin{bmatrix} 2 \\ 2 \\ 4 \\ 2 \\ 2 \\ 8 \\ 6 \end{bmatrix} < \begin{bmatrix} 2 \\ 3 \\ 5 \\ 6 \\ 8 \\ 9 \end{bmatrix}$$

$$\begin{bmatrix} 1 \\ 2 \\ 4 \\ 2 \\ 5 \\ 6 \end{bmatrix} = \begin{bmatrix} 3 \\ 5 \\ 3 \\ 1 \\ 8 \\ 9 \end{bmatrix}$$

$$\max \left( \begin{bmatrix} 1 \\ 2 \\ 4 \\ 2 \\ 5 \\ 6 \end{bmatrix}, \begin{bmatrix} 3 \\ 5 \\ 3 \\ 1 \\ 8 \\ 9 \end{bmatrix} \right) = \begin{bmatrix} 3 \\ 5 \\ 4 \\ 2 \\ 8 \\ 9 \end{bmatrix}$$

# Vector Clocks – Example



# Relation Between the Time Stamps

- >  $R(e)$  exact real time of the event  $e$
- >  $L(e)$  Lamport-time of  $e$
- >  $V(e)$  Vector-time of  $e$

$$\begin{array}{ccc} R(a) < R(b) & \Leftarrow & V(a) < V(b) \\ \Uparrow & \nearrow & \Downarrow \\ a \rightarrow b & \Rightarrow & L(a) < L(b) \end{array}$$

# Exemplary Applications of Logical Clocks

# Application of Vector Clocks: Causal Multicast

- > Each message is delivered to all processes
- > Messages must be delivered in an order satisfying causality
- > Update policy of the vector clocks is changed such that  $P_i$  only increments  $V_i[i]$  if it sends a message
- > **Delivery condition**
  - > A message with time stamp  $T$  sent by  $P_i$  is not delivered to  $P_j$  before  $T$  fulfills the following conditions:

$$\underbrace{T[i] = V_j[i] + 1}_{\text{Next message awaited from } P_i} \wedge \underbrace{\forall k \neq i: T[k] \leq V_j[k]}_{\text{No message from arbitrary } P_k \text{ is missing}}$$


- > 1<sup>st</sup> condition violated  $\rightarrow$  a message from  $P_i$  is missing
- > 2<sup>nd</sup> condition violated  $\rightarrow$  a message from  $P_k$  is missing that  $P_i$  had already received at the time it sent the message received yet

# Application of Vector Clocks: Causal Multicast

- > The maximum of  $V_j$  und  $T$  is derived after the delivery of the message
- > Due to the second delivery condition:
$$V_j[k] \geq T[k] \text{ for all } k \neq i$$
- > This means that except for the component  $V_j[i]$ ,  $V_j$  only contains components that are larger or equal to those in  $T$
- > Due to the first delivery condition, for  $V_j[i]$  holds:
$$V_j[i] = T[i] - 1$$
- > Therefore, calculating the maximum is not necessary
- > Instead, it is sufficient to increase  $V_j[i]$  by one at delivery



# Causal Multicast – Example



# Exemplary Exam Questions

1. Explain the external clock synchronization algorithm!
2. What is Lamport's clock condition?
3. How the Happened-Before-Relation is defined?
4. Describe the functioning of Lamport clocks!
5. Can events be arranged by Lamport clocks in such a way that causally related events are properly ordered?
6. Explain the functionality of vector clocks!
7. What additional opportunities vector clocks can provide compared to Lamport clocks?
8. Describe how a causal multicast can be realized using vector clocks!

# Literature

1. A. S. Tanenbaum and M. van Steen. *Distributed Systems: Principles and Paradigms*. Prentice Hall, 2002. Chapter 5.1 + 5.2
2. G. Coulouris, J. Dollimore, and T. Kindberg. *Distributed Systems: Concepts and Design*. Addison-Wesley, 4th edition, 2005. Chapter 11.1 – 11.4
3. N. Lynch. *Distributed Algorithms*. Morgan Kaufmann, 1996. Chapter 18
4. F. Mattern. *Verteilte Basisalgorithmen*. Springer-Verlag, 1989. Kapitel 5: Virtuelle Zeit in verteilten Systemen
5. F. Cristian. *Probabilistic clock synchronization*. Distributed Computing, 3(3):146--158, September 1989.

# Literature

6. D. L. Mills. *Network Time Protocol (Version 3)*. Internet Request for Comments RFC 1305, March 1992.
7. L. Lamport. *Time, Clocks, and the Ordering of Events in a Distributed Environment*. Communications of the ACM, 21:558--564, July 1978.
8. F. Mattern. *Virtual Time and Global States of Distributed Systems*. In C. M. et al., editor, Proc. Workshop on Parallel and Distributed Algorithms, pages 215--226, North-Holland / Elsevier, 1989. (Reprinted in: Z. Yang, T. A. Marsland (Eds.), "Global States and Time in Distributed Systems", IEEE, 1994, pp. 123-133.).
9. Michel Raynal. *About logical clocks for distributed systems*. SIGOPS Oper. Syst. Rev. 26, pp. 41-48, 1992.

# Thank you for your kind attention!

**Univ.-Prof. Dr.-Ing. habil. Gero Mühl**

`gero.muehl@uni-rostock.de`

`http://wwwava.informatik.uni-rostock.de`